


BS 6207 ASSIGNMENT 1 - Tian Siqu, G2101015G

Github link(codes,output files and a more detailed report):

<https://github.com/SiqiT/Assignment1>

Question 1

Here we set $k=10$, $d=10$ (10 hidden layers, each layer has 21 nodes)



Input layer: $X.shape(1,10) \rightarrow (batch_size=1, 10 \text{ nodes})$

Output layer: $y.shape(1,1) \rightarrow (batch_size=1, 1 \text{ node})$

```
X = Variable(torch.rand(batch_size, d), requires_grad=True)
y = torch.tensor([[torch.sum(X**2)/d]])
```

Activation function *ReLU* and its derivative:

```
def ReLU_d(x):
    return np.where(x < 0, 0, 1)
def ReLU(x):
    return np.where(x < 0, 0, x)
```

Loss function: $loss = (y_pred - y)**2$

(1) Forward Propagation

- Autograd

I use the *nn_module* class to define the network structure and achieve the forward propagation for one time. '*nn.Linear*' is used for full_connected layer, '*nn.ReLU*' is used for the activation function.

- My_grad

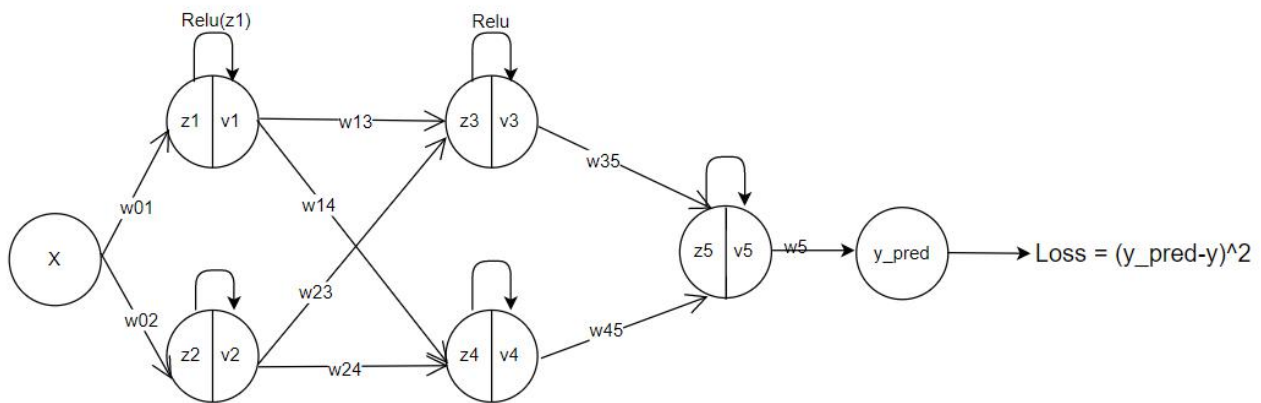
Based on the same original weight and bias in the autograd, I use *np.dot* to achieve forward propagation in the function feedforward and save the node value per layer in a tensor variable.

(2) Backward Propagation

- Autograd

loss.backward() can compute $dloss/dx$ for every parameter, then use *param.grad* to get grad for each variable.

- My_grad



Activation function:

$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases} \quad f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$

Loss:

$$\frac{dL}{dy_{\text{pred}}} = 2 * (y_{\text{pred}} - y)$$

Output layer:

$$y_{\text{pred}} = v_5 * w_5 + b_5 \quad v_5 = \delta(z_5)$$

$$\frac{\partial L}{\partial w_5} = \frac{\partial L}{\partial y_{\text{pred}}} \frac{\partial y_{\text{pred}}}{\partial w_5} = dL * \delta'(z_5) * v_5 = 2(y_{\text{pred}} - y) * \delta'(z_5) * v_5$$

$$\frac{\partial L}{\partial b_5} = \frac{\partial L}{\partial y_{\text{pred}}} \frac{\partial y_{\text{pred}}}{\partial b_5} = dL * \delta'(z_5) = 2(y_{\text{pred}} - y) * \delta'(z_5)$$

The same, Hidden layer:

$$\frac{\partial L}{\partial w_4} = \frac{\partial L}{\partial z_5} \frac{\partial z_5}{\partial w_4} = \frac{\partial L}{\partial z_5} * \delta'(z_4) * w_{45} * v_4$$

$$\frac{\partial L}{\partial b_4} = \frac{\partial L}{\partial z_5} \frac{\partial z_5}{\partial b_4} = \frac{\partial L}{\partial z_5} * \delta'(z_4) * w_{45}$$

Input layer: change v to x ;

Finally, Compare the two files torch_autograd.dat and my_autograd.dat and show that they give the same values for up to 5 significant numbers:

```
diff_w=0
for i in range(len(d_w)):
    diff_w = diff_w + np.absolute(sum(sum(d_w[i]-auto_w[i])))
```

```
diff_w
```

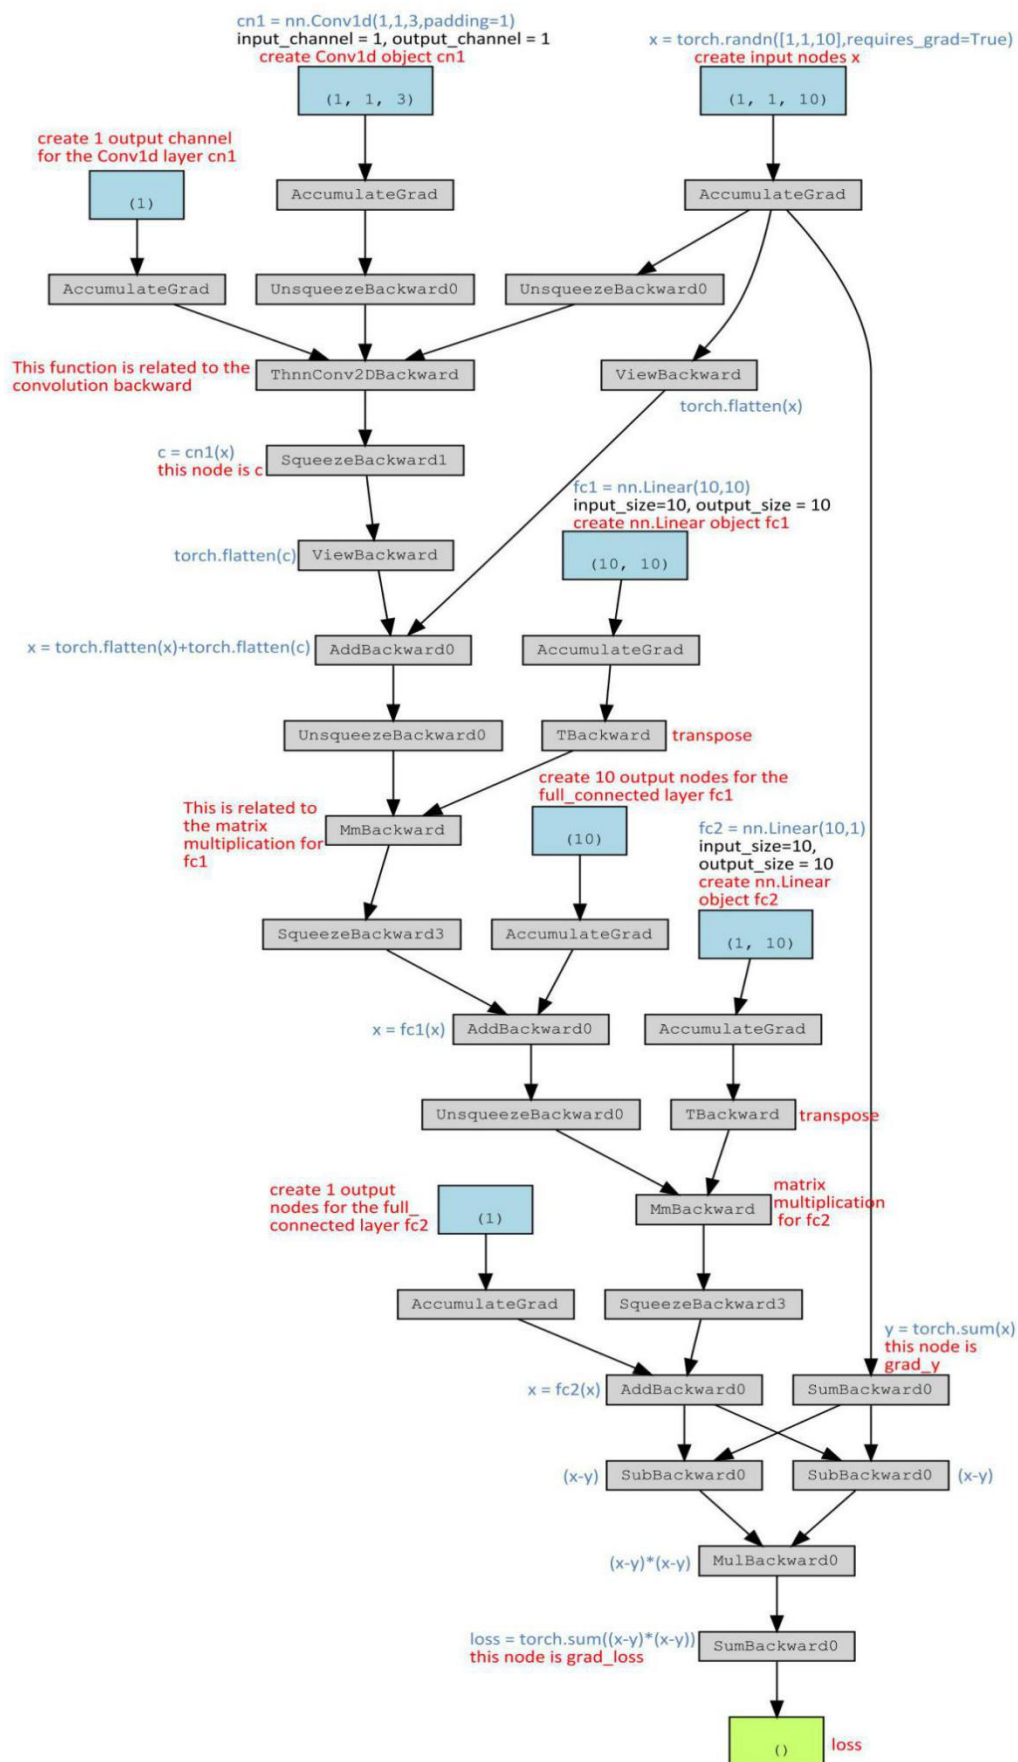
```
0.0
```

```
diff_b=0
for i in range(len(d_z)):
    diff_b = diff_b + np.absolute(sum(sum(d_z[i]-auto_b[i])))
diff_b
```

```
0.0
```

The gradient values from those 2 ways are the same.

Question 2



This is the computational graph for the above scripts, and there are explanations for nodes in the graph above.

For different types of tensor variables, the autograd backward functions for calculating gradients would be different:

```
torch.sum() -> grad_fn=<SumBackward>  
nn.Conv1d(x) -> grad_fn=<SqueezeBackward>  
torch.flatten() -> grad_fn=<ViewBackward>  
a+b -> grad_fn=<AddBackward>  
a-b -> grad_fn=<SubBackward>  
a*b -> grad_fn=<MulBackward>  
nn.Linear(x) -> grad_fn=<AddBackward>
```

For other nodes that are not marked in the graph:

'AccumulateGrad' represents leaf nodes, also as known as the end point of the computational graph for BP. It accumulates all backward gradient information for the leaf nodes.

'UnsqueezeBackward', *'SqueezeBackward'* are working with the tensor variable's dimension and shape, to add a dimension or remove a dimension for a variable.